

NBA Scoring Analysis

Aidan Berry

November 12, 2018

Overall Question

Can we accurately predict the number of points a player will score over the course of a season using their demographics, attributes, position, team, and performance data?

Step 1: Data Preprocessing and EDA

Here I create the normalize function that is used to standardize the data on a uniform scale as well as the RMSE function which calculates the Root Mean Square Error for a regression model and will be used to test the accuracy of the models created. I also create a function for one hot encoding categorical variables in the dataset.

```
normalize <- function(x) {
  num <- x - mean(x)
  denom <- sd(x)
  return (num/denom)
}

rmse <- function(y,yhat) {
  num <- sum((y - yhat)^2)
  denom <- length(y)
  return(sqrt(num/denom))
}

onehotencoder <- function(df_orig) {
  df<-cbind(df_orig)
  df_clmtyp<-data.frame(clmtyp=sapply(df,class))
  df_col_typ<-data.frame(clnmn=colnames(df),clmtyp=df_clmtyp$clmtyp)
  for (rownm in 1:nrow(df_col_typ)) {
    if (df_col_typ[rownm,"clmtyp"]=="factor") {
      clmn_obj<-df[tostring(df_col_typ[rownm,"clnmn"])]
      dummy_matx<-data.frame(model.matrix( ~.-1, data = clmn_obj))
      dummy_matx<-dummy_matx[,c(1,3:ncol(dummy_matx))]
      df[tostring(df_col_typ[rownm,"clnmn"])]<-NULL
      df<-cbind(df,dummy_matx)
      df[tostring(df_col_typ[rownm,"clnmn"])]<-NULL
    }
  }
  return(df)
}
```

I now read in the dataset from an exported SQL query that joins fields from 3 of the tables. This dataset still needs to be adjusted however, to get the final cleaned dataset for analysis. We will use dplyr for this.

```
dataset <- read.csv('data/regressorPreCleanedData.csv')
allStar <- read.csv('data/nba_all_star_games.csv')
```

The only column with null values is the college column with 2163 null values. To impute the null values for the college variable, I set all of the null values to the string “No College”, which will act as a new category in this column. To do this, I had to first change the type of the College column to character, impute the missing values, and then convert the datatype back to factor (categorical). A player not going to college is still valuable data to include in the model because there could be a correlation between the players that didn’t go to college and points scored in a season.


```
dataset$College <- as.character(dataset$College)
dataset$College[dataset$College==""] <- "No College"
dataset$College <- as.factor(dataset$College)
```

Here we need to determine if each player in the dataset was an All Star at any point in their career. To do this I created a vector that contains whether each player in the dataset was contained in the allStar dataset by signifying true or false. I then used the vectorized functionality of R to go through the dataset and change the All_Star variable to 1 if the respective value in the allstarstatus array is true.

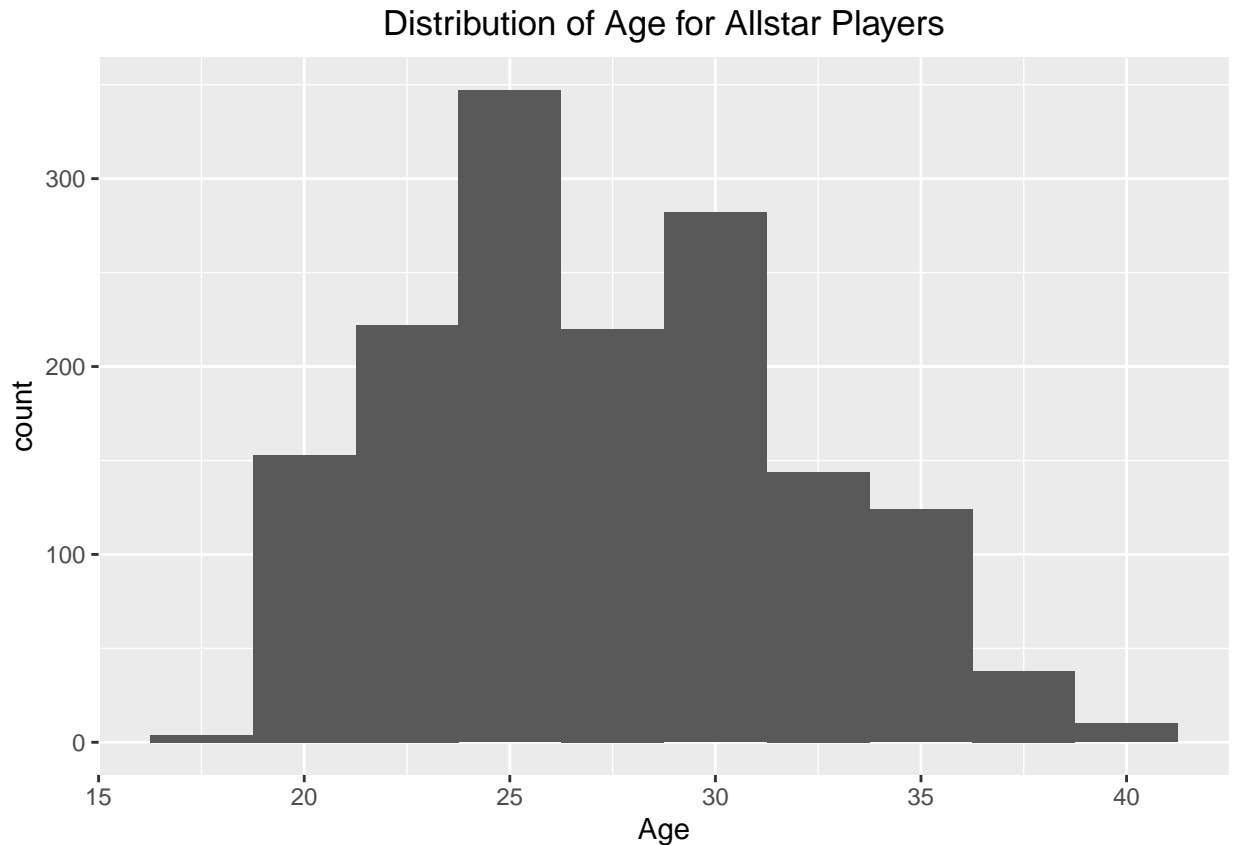
```
allstarstatus <- dataset$Player_Name %in% allStar$player
dataset$All_Star[allstarstatus==TRUE] <- 1
```

This is now the final version of the dataset that will be used. A description of each variable is described below.

- Points: Total number of points the player scored during that season
- Team: 3 letter abbreviation for the team that the player is on
- height: Height of the player in cm
- weight: Weight of the player in kg
- Position: The abbreviation for the position the player plays on the court
- Age: The age of the player in years
- All_Star: 1 if the player was ever on an all-star team, 0 if not



Warning: Use of `allStarAge\$Age` is discouraged. Use `Age` instead.




It appears that the SG-PF position and PG-SF positions score the most points on average per season. This makes sense because they are attacking positions that have a lot of time on the ball and are taking the most shots.


To get the final dataset that is going to be used in the models, all of the variables must be numeric, which means we have to encode the 2 categorical variables and drop unnecessary ones. Here I dropped `Player_Name` and `College`, because the name doesn't provide any relevant data and there are too many different factors for college that if it was oneHotEncoded, there would be a lot of noise and too many dimensions in the dataset for the models to make any sense of it.


```
dataset <- dataset[-c(1,7)]
encodedDataset <- onehotencoder(dataset)
```

Here I visualized the relationship between the variables using a correlation plot. To build this I dropped the `Team` and `Position` variables because there is too many levels to make sense of in the chart, and it makes the chart easier to read without them. From this we can see there is a very strong relationship between height and weight and a small correlation between `all_star` status and `Points scored`, which both make sense in this case.

```
## Warning: package 'ggcorrplot' was built under R version 3.6.3
```







From these plots we can see that most of the all-stars are in the upper-middle height range from about 185 cm - 220 cm. The low end and very high end of the height range have extremely low numbers of all-stars. The number of points scored follows a similar pattern, with the most average season points being scored by players in the upper-middle height range, with dips in points near both height extremes.

The age of players also plays an important role in the number of points scored by players as well as their all-star status. The age range of 21-39 seems to produce the most all-stars. The optimal player age for scoring the most average season points appears to be around the 23-29 age range. The points scored starts low and peaks in this 23-29 range, and then slowly decreases from there on out as the player gets older.

Step 2: Building the Models

In order to build the model, we must first split the data into a training set and a test set. I use the split method from the caTools library to do this in one line with a train/test ratio of 75/25 since we have a pretty good sized dataset and want the test set to be as large as possible for validation. I then also created a scaled version of the train and test sets that are used in some models that use a euclidean distance algorithm, this way all the data is on the same scale.

```
set.seed(123)
split <- sample.split(encodedDataset$Points, SplitRatio = 0.75)
training_set <- subset(encodedDataset, split == TRUE)
test_set <- subset(encodedDataset, split == FALSE)
scaled_training_set =training_set
scaled_test_set=test_set
scaled_training_set[2:4] = scale(training_set[2:4])
scaled_test_set[2:4] = scale(test_set[2:4])
```

Multiple Linear Regression

The first model I tried to build for this dataset is the multiple linear regression model. I created the regressor and then predicted the values of the test set using this regressor. The data doesn't need to be scaled for linear regression, so I just used the original training and test sets.

```
set.seed(123)
regressor_MLR = lm(formula = Points ~ ., data = training_set)
summary(regressor_MLR)
```

```
##
## Call:
## lm(formula = Points ~ ., data = training_set)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-1141.3	-360.5	-121.5	262.7	3443.1

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-560.7151	185.3858	-3.025	0.002493	**
height	4.7129	0.8087	5.828	5.72e-09	***
weight	-1.9143	0.5551	-3.449	0.000565	***
Age	1.6177	0.9530	1.698	0.089609	.
All_Star	533.5736	14.7691	36.128	< 2e-16	***
TeamAND	38.8764	148.5392	0.262	0.793537	
TeamBAL	179.3321	49.9466	3.590	0.000331	***
TeamBLB	-46.2017	65.0888	-0.710	0.477822	
TeamBOS	49.1382	26.7228	1.839	0.065959	.
TeamBRK	-180.7105	60.9671	-2.964	0.003040	**
TeamBUF	83.8548	53.0676	1.580	0.114090	
TeamCAP	-35.6046	161.5784	-0.220	0.825597	
TeamCHA	-118.1194	46.9135	-2.518	0.011817	*
TeamCHH	-42.4861	42.4075	-1.002	0.316429	
TeamCHI	-12.3060	28.5094	-0.432	0.666002	
TeamCHO	-82.6903	76.0732	-1.087	0.277059	
TeamCHP	231.1336	161.5727	1.431	0.152584	
TeamCHS	35.8977	148.7117	0.241	0.809255	
TeamCHZ	100.6865	130.4806	0.772	0.440327	
TeamCIN	251.1213	43.5967	5.760	8.55e-09	***
TeamCLE	-39.0598	28.5239	-1.369	0.170901	
TeamDAL	-76.0493	30.4069	-2.501	0.012391	*
TeamDEN	25.9999	29.7790	0.873	0.382622	
TeamDET	33.3179	27.3382	1.219	0.222962	
TeamDNN	111.1302	143.3652	0.775	0.438258	
TeamFTW	8.5226	59.3995	0.143	0.885913	
TeamGSW	-6.9987	28.6881	-0.244	0.807266	
TeamHOU	8.6936	28.8574	0.301	0.763219	
TeamIND	1.3888	30.1715	0.046	0.963286	
TeamINO	22.1035	81.8182	0.270	0.787045	
TeamKCK	110.7016	48.6490	2.276	0.022887	*
TeamKCO	135.3886	84.9204	1.594	0.110886	
TeamLAC	-91.0691	31.0097	-2.937	0.003320	**
TeamLAL	104.9419	28.0566	3.740	0.000184	***

## TeamMEM	-143.0968	38.8434	-3.684	0.000230	***
## TeamMIA	-86.4114	32.8131	-2.633	0.008460	**
## TeamMIL	8.7003	28.3994	0.306	0.759338	
## TeamMIN	-67.9701	33.8010	-2.011	0.044352	*
## TeamMLH	-151.3257	67.5873	-2.239	0.025170	*
## TeamMNL	100.7422	50.7817	1.984	0.047290	*
## TeamNJN	-72.3264	30.2280	-2.393	0.016735	*
## TeamNOH	-138.0742	48.6892	-2.836	0.004576	**
## TeamNOJ	42.2836	68.4937	0.617	0.537022	
## TeamNOK	-143.7665	102.3421	-1.405	0.160109	
## TeamNOP	-157.5617	63.4175	-2.485	0.012982	*
## TeamNYK	23.7473	26.6015	0.893	0.372026	
## TeamNYN	-135.5619	134.9178	-1.005	0.315019	
## TeamOKC	-55.6858	49.9507	-1.115	0.264945	
## TeamORL	-45.8033	33.5532	-1.365	0.172241	
## TeamPHI	17.7571	27.9855	0.635	0.525754	
## TeamPHO	26.6412	28.5616	0.933	0.350955	
## TeamPHW	135.4549	49.2234	2.752	0.005932	**
## TeamPOR	19.6365	29.0036	0.677	0.498392	
## TeamROC	55.0857	63.5101	0.867	0.385760	
## TeamSAC	-45.1278	31.6841	-1.424	0.154376	
## TeamSAS	-40.8435	29.5156	-1.384	0.166439	
## TeamSDC	69.3029	63.4290	1.093	0.274582	
## TeamSDR	268.6761	79.6228	3.374	0.000741	***
## TeamSEA	72.0820	30.5049	2.363	0.018140	*
## TeamSFW	242.1999	54.5720	4.438	9.13e-06	***
## TeamSHE	70.9726	144.6030	0.491	0.623567	
## TeamSTB	183.0678	185.9473	0.985	0.324876	
## TeamSTL	215.4633	47.8741	4.501	6.82e-06	***
## TeamSYR	154.8073	48.6664	3.181	0.001470	**
## TeamTOR	-137.3903	34.2033	-4.017	5.92e-05	***
## TeamTOT	-106.9614	23.9308	-4.470	7.88e-06	***
## TeamTRI	-78.0608	91.0928	-0.857	0.391492	
## TeamUTA	4.4904	30.7473	0.146	0.883890	
## TeamVAN	-120.3713	61.6956	-1.951	0.051067	.
## TeamWAS	-135.5263	36.2507	-3.739	0.000186	***
## TeamWAT	-130.5845	142.6786	-0.915	0.360080	
## TeamWSB	33.8533	35.7881	0.946	0.344193	
## TeamWSC	-39.1471	109.0927	-0.359	0.719717	
## PositionC	190.1067	113.2352	1.679	0.093196	.
## PositionC.PF	172.9163	153.0826	1.130	0.258676	
## PositionC.SF	105.0891	358.5156	0.293	0.769432	
## PositionF	-149.2094	124.6326	-1.197	0.231247	
## PositionF.C	65.4375	130.4134	0.502	0.615836	
## PositionF.G	-72.9566	136.0309	-0.536	0.591742	
## PositionG	-60.0018	123.4561	-0.486	0.626961	
## PositionG.F	50.6090	130.7761	0.387	0.698769	
## PositionPF	250.9197	113.1604	2.217	0.026610	*
## PositionPF.C	119.4372	153.0693	0.780	0.435236	
## PositionPF.SF	285.0312	162.8290	1.750	0.080050	.
## PositionPG	298.1968	113.7526	2.621	0.008763	**
## PositionPG.SF	864.6555	494.1498	1.750	0.080173	.
## PositionPG.SG	287.8840	154.9380	1.858	0.063177	.
## PositionSF	323.3949	113.1982	2.857	0.004283	**

```
## PositionSF.PF 291.4697 165.4397 1.762 0.078122 .
## PositionSF.PG 413.3349 494.1680 0.836 0.402927
## PositionSF.SG 348.6578 154.7068 2.254 0.024229 *
## PositionSG 340.4957 113.3569 3.004 0.002670 **
## PositionSG.PF 850.2253 300.1392 2.833 0.004620 **
## PositionSG.PG 280.0928 153.4021 1.826 0.067886 .
## PositionSG.SF 380.1852 171.6105 2.215 0.026746 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 480.9 on 17811 degrees of freedom
## Multiple R-squared:  0.1049, Adjusted R-squared:  0.1002
## F-statistic: 22.21 on 94 and 17811 DF,  p-value: < 2.2e-16
```

```
y_pred_MLR = predict(regressor_MLR, newdata = test_set)
rmse_MLR <- rmse(test_set$Points, y_pred_MLR)
```

```
# look at the vif values of the linear regression to see if there is multicollinearity
vif(regressor_MLR)
```

```
##      height      weight      Age      All_Star      TeamAND
##      4.441047      3.452289      1.029005      1.051111      1.048846
##      TeamBAL      TeamBLB      TeamBOS      TeamBRK      TeamBUF
##      1.189974      1.186453      2.213280      1.120705      1.162782
##      TeamCAP      TeamCHA      TeamCHH      TeamCHI      TeamCHO
##      1.015533      1.218847      1.278948      1.929691      1.073474
##      TeamCHP      TeamCHS      TeamCHZ      TeamCIN      TeamCLE
##      1.015462      1.051283      1.029872      1.262893      1.925068
##      TeamDAL      TeamDEN      TeamDET      TeamDNN      TeamFTW
##      1.742508      1.787650      2.093246      1.065813      1.154612
##      TeamGSW      TeamHOU      TeamIND      TeamINO      TeamKCK
##      1.903915      1.885870      1.760479      1.097652      1.199668
##      TeamKCO      TeamLAC      TeamLAL      TeamMEM      TeamMIA
##      1.058232      1.689593      1.989768      1.354046      1.575320
##      TeamMIL      TeamMIN      TeamMLH      TeamMNL      TeamNJN
##      1.940953      1.524088      1.102711      1.197055      1.748335
##      TeamNOH      TeamNOJ      TeamNOK      TeamNOP      TeamNYK
##      1.201655      1.092162      1.040355      1.109038      2.240857
##      TeamNYN      TeamOKC      TeamORL      TeamPHI      TeamPHO
##      1.022515      1.190171      1.534692      1.995475      1.920235
##      TeamPHW      TeamPOR      TeamROC      TeamSAC      TeamSAS
##      1.207487      1.870825      1.112280      1.639491      1.823838
##      TeamSDC      TeamSDR      TeamSEA      TeamSFW      TeamSHE
##      1.109442      1.066833      1.726987      1.153182      1.084298
##      TeamSTB      TeamSTL      TeamSYR      TeamTOR      TeamTOT
##      1.046193      1.210641      1.210635      1.506848      3.487678
##      TeamTRI      TeamUTA      TeamVAN      TeamWAS      TeamWAT
##      1.146160      1.707866      1.114984      1.428388      1.055628
##      TeamWSB      TeamWSC      PositionC      PositionC.PF      PositionC.SF
##      1.440569      1.079456      155.077705      2.226609      1.111478
##      PositionF      PositionF.C      PositionF.G      PositionG      PositionG.F
##      5.416048      3.886304      3.273156      6.032254      3.907948
##      PositionPF      PositionPF.C      PositionPF.SF      PositionPG      PositionPG.SF
```

```
##      160.516627      2.226223      1.947167      153.283677      1.055836
## PositionPG.SG      PositionSF PositionSF.PF PositionSF.PG PositionSF.SG
##      2.177355      152.550948      1.891971      1.055914      2.170860
##      PositionSG PositionSG.PF PositionSG.PG PositionSG.SF
##      157.273811      1.168416      2.235914      1.781475
```

```
# print the variables that have a vif value higher than 10 (multicollinearity problem)
which(vif(regressor_MLR)>10)
```

```
## PositionC PositionPF PositionPG PositionSF PositionSG
##      73      81      84      87      91
```

```
set.seed(123)
regressor_MLR2 = lm(formula = Points ~ . - PositionPF, data = training_set)
summary(regressor_MLR2)
```

```
##
## Call:
## lm(formula = Points ~ . - PositionPF, data = training_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1141.8   -361.0   -121.4    263.0   3447.3
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -308.1241   146.2743  -2.106  0.035177 *
## height         4.6835     0.8087   5.791  7.10e-09 ***
## weight        -1.8855     0.5550  -3.397  0.000682 ***
## Age           1.6329     0.9531   1.713  0.086666 .
## All_Star      533.6402    14.7707  36.128 < 2e-16 ***
## TeamAND       35.3355    148.5470   0.238  0.811981
## TeamBAL      179.2347    49.9521   3.588  0.000334 ***
## TeamBLB      -52.2397    65.0390  -0.803  0.421866
## TeamBOS       47.6815    26.7176   1.785  0.074336 .
## TeamBRK     -180.7713    60.9738  -2.965  0.003033 **
## TeamBUF       83.9146    53.0734   1.581  0.113872
## TeamCAP      -35.6714   161.5961  -0.221  0.825294
## TeamCHA     -118.2015    46.9187  -2.519  0.011768 *
## TeamCHH      -42.5327    42.4122  -1.003  0.315951
## TeamCHI      -12.2725    28.5125  -0.430  0.666892
## TeamCHO      -82.7304    76.0816  -1.087  0.276879
## TeamCHP      231.1976   161.5905   1.431  0.152516
## TeamCHS      -11.3770   147.1916  -0.077  0.938391
## TeamCHZ       99.1505   130.4931   0.760  0.447376
## TeamCIN      250.9706    43.6014   5.756  8.75e-09 ***
## TeamCLE      -39.5539    28.5262  -1.387  0.165587
## TeamDAL      -76.0976    30.4102  -2.502  0.012345 *
## TeamDEN       25.9840    29.7822   0.872  0.382965
## TeamDET       33.3188    27.3412   1.219  0.223001
## TeamDNN      107.6356   143.3722   0.751  0.452818
## TeamFTW        4.0403    59.3717   0.068  0.945746
## TeamGSW      -6.9963    28.6913  -0.244  0.807351
```

## TeamHOU	8.7004	28.8605	0.301	0.763065	
## TeamIND	1.3785	30.1748	0.046	0.963562	
## TeamINO	13.9230	81.7440	0.170	0.864757	
## TeamKCK	110.7879	48.6543	2.277	0.022795	*
## TeamKCO	135.4859	84.9297	1.595	0.110669	
## TeamLAC	-91.1186	31.0131	-2.938	0.003307	**
## TeamLAL	104.9077	28.0597	3.739	0.000186	***
## TeamMEM	-143.1616	38.8476	-3.685	0.000229	***
## TeamMIA	-86.4824	32.8167	-2.635	0.008413	**
## TeamMIL	8.6805	28.4025	0.306	0.759896	
## TeamMIN	-68.0146	33.8048	-2.012	0.044237	*
## TeamMLH	-156.3346	67.5570	-2.314	0.020673	*
## TeamMNL	100.0844	50.7864	1.971	0.048775	*
## TeamNJN	-72.2624	30.2313	-2.390	0.016844	*
## TeamNOH	-138.2465	48.6945	-2.839	0.004530	**
## TeamNOJ	42.3290	68.5013	0.618	0.536630	
## TeamNOK	-143.6439	102.3533	-1.403	0.160511	
## TeamNOP	-157.6888	63.4244	-2.486	0.012919	*
## TeamNYK	22.6383	26.5997	0.851	0.394740	
## TeamNYN	-135.5162	134.9326	-1.004	0.315236	
## TeamOKC	-55.8222	49.9562	-1.117	0.263828	
## TeamORL	-45.8205	33.5569	-1.365	0.172128	
## TeamPHI	17.7432	27.9886	0.634	0.526125	
## TeamPHO	26.6670	28.5647	0.934	0.350542	
## TeamPHW	130.7063	49.1822	2.658	0.007877	**
## TeamPOR	19.6401	29.0067	0.677	0.498359	
## TeamROC	54.7507	63.5169	0.862	0.388707	
## TeamSAC	-45.1658	31.6876	-1.425	0.154075	
## TeamSAS	-40.8132	29.5188	-1.383	0.166800	
## TeamSDC	69.3626	63.4360	1.093	0.274221	
## TeamSDR	268.6499	79.6316	3.374	0.000743	***
## TeamSEA	72.0682	30.5082	2.362	0.018175	*
## TeamSFW	242.0277	54.5780	4.435	9.28e-06	***
## TeamSHE	66.2804	144.6035	0.458	0.646700	
## TeamSTB	144.5473	185.1544	0.781	0.434998	
## TeamSTL	215.3001	47.8793	4.497	6.94e-06	***
## TeamSYR	154.3701	48.6714	3.172	0.001518	**
## TeamTOR	-137.4567	34.2070	-4.018	5.88e-05	***
## TeamTOT	-107.0432	23.9334	-4.473	7.78e-06	***
## TeamTRI	-96.9329	90.7043	-1.069	0.285233	
## TeamUTA	4.4370	30.7506	0.144	0.885274	
## TeamVAN	-120.3916	61.7024	-1.951	0.051053	.
## TeamWAS	-135.5996	36.2546	-3.740	0.000184	***
## TeamWAT	-134.1504	142.6852	-0.940	0.347136	
## TeamWSB	33.8972	35.7920	0.947	0.343621	
## TeamWSC	-42.1110	109.0965	-0.386	0.699503	
## PositionC	-59.5600	12.0223	-4.954	7.33e-07	***
## PositionC.PF	-76.9951	103.6053	-0.743	0.457396	
## PositionC.SF	-144.5671	340.4148	-0.425	0.671076	
## PositionF	-395.7511	56.3173	-7.027	2.18e-12	***
## PositionF.C	-179.1648	69.5710	-2.575	0.010024	*
## PositionF.G	-319.6323	78.2940	-4.082	4.48e-05	***
## PositionG	-305.4176	54.7028	-5.583	2.40e-08	***
## PositionG.F	-192.1964	71.5055	-2.688	0.007198	**

```
## PositionPF.C -130.4076 103.6184 -1.259 0.208214
## PositionPF.SF 35.3004 117.6099 0.300 0.764067
## PositionPG 48.5597 16.2789 2.983 0.002858 **
## PositionPG.SF 614.9081 481.1954 1.278 0.201310
## PositionPG.SG 38.0787 106.3790 0.358 0.720382
## PositionSF 73.8061 12.0024 6.149 7.95e-10 ***
## PositionSF.PF 41.6152 121.1475 0.344 0.731220
## PositionSF.PG 163.6593 481.2216 0.340 0.733792
## PositionSF.SG 98.8833 106.0566 0.932 0.351161
## PositionSG 90.9355 13.5252 6.723 1.83e-11 ***
## PositionSG.PF 600.4719 278.2338 2.158 0.030929 *
## PositionSG.PG 30.2566 104.1137 0.291 0.771352
## PositionSG.SF 130.3670 129.4633 1.007 0.313958
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 480.9 on 17812 degrees of freedom
## Multiple R-squared: 0.1047, Adjusted R-squared: 0.1
## F-statistic: 22.39 on 93 and 17812 DF, p-value: < 2.2e-16
```

```
y_pred_MLR2 = predict(regressor_MLR2, newdata = test_set[-1])
rmse_MLR2 <- rmse(test_set$Points, y_pred_MLR2)
```

```
# this fixed the multicollinearity problem but the regression fit is still not great
which(vif(regressor_MLR2)>10)
```

```
## named integer(0)
```

```
# perform stepwise selection to pick the variables in the model systematically
set.seed(123)
train_control = trainControl(method="cv", number = 10)
step_wise_model = train(Points~., data = training_set, method = "leapSeq",
                        tuneGrid=data.frame(nvmax=1:95), trControl = train_control)
```

```
## Warning in leaps.setup(x, y, wt = weights, nbest = nbest, nvmax = nvmax, : 1
## linear dependencies found
```

```
## Reordering variables and trying again:
```

```
## Warning: predictions failed for Fold02: nvmax=95 Error in method$predict(modelFit = modelFit, newdata = newdata,
## Some values of 'nvmax' are not in the model sequence.
```

```
## Warning: 1 linear dependencies found
```

```
## Reordering variables and trying again:
```

```
## Warning: predictions failed for Fold10: nvmax=95 Error in method$predict(modelFit = modelFit, newdata = newdata,
## Some values of 'nvmax' are not in the model sequence.
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.
```

```
# the best model contained 81 variables
step_wise_model$results
```

##	nvmax	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 1	1	492.5847	0.06352857	389.0463	7.143910	0.008904095	4.942410
## 2	2	491.7208	0.06672220	388.3222	7.088769	0.008543061	4.956679
## 3	3	490.5960	0.07087496	386.6737	6.897200	0.007146835	4.917954
## 4	4	489.8972	0.07354328	385.7532	7.120677	0.008061656	5.409516
## 5	5	489.2403	0.07597686	384.9611	7.134440	0.007892107	5.490038
## 6	6	489.3207	0.07561339	385.2398	7.148749	0.007645635	5.457132
## 7	7	489.7376	0.07415941	385.9948	7.389400	0.011941553	5.576446
## 8	8	488.6554	0.07817350	384.8143	7.353211	0.008669677	5.389487
## 9	9	488.6990	0.07819111	385.2012	7.052943	0.007229905	5.218492
## 10	10	488.3381	0.07931297	384.6365	6.711564	0.007332499	5.281028
## 11	11	488.6221	0.07826411	385.2850	7.843621	0.008342485	6.526024
## 12	12	488.9506	0.07711223	385.5851	7.048728	0.008386094	5.084662
## 13	13	488.4990	0.07887248	385.1671	7.464287	0.007045598	5.810736
## 14	14	487.5272	0.08245827	384.3869	7.494005	0.009102786	5.859294
## 15	15	487.7242	0.08162322	384.5587	7.328101	0.008619866	5.467925
## 16	16	487.6946	0.08191508	384.7984	6.928362	0.006077059	4.934461
## 17	17	486.9496	0.08457120	384.0155	7.006284	0.007433052	5.423698
## 18	18	487.7161	0.08181544	384.8429	8.437272	0.011026762	6.823368
## 19	19	487.4296	0.08290342	384.6637	5.922558	0.006481139	4.469669
## 20	20	486.5401	0.08613758	383.6505	7.146028	0.008222849	5.569326
## 21	21	486.9716	0.08439405	383.9966	6.151242	0.007707214	5.079312
## 22	22	486.9930	0.08443266	384.2062	8.339564	0.011847740	6.569507
## 23	23	486.5475	0.08627297	383.9411	6.810181	0.005690683	5.029694
## 24	24	486.4397	0.08663027	383.7897	7.232517	0.006554616	5.859820
## 25	25	486.5721	0.08604435	383.8303	8.517333	0.012729805	6.739643
## 26	26	486.0650	0.08797212	383.2594	7.094368	0.007212542	5.516875
## 27	27	485.8098	0.08894053	383.0747	7.102851	0.007767159	5.541129
## 28	28	486.8167	0.08512197	383.9216	7.286400	0.012673942	5.444084
## 29	29	486.0663	0.08790965	383.2648	7.034939	0.009690190	5.585154
## 30	30	485.3210	0.09073860	382.5827	6.935527	0.007592797	5.449009
## 31	31	485.1965	0.09118475	382.4607	6.908309	0.007337291	5.466141
## 32	32	485.1236	0.09146449	382.4005	6.835345	0.007503115	5.456363
## 33	33	485.3814	0.09040336	382.4678	6.244451	0.008714691	5.316571
## 34	34	485.7592	0.08918899	382.9053	8.587590	0.013811108	6.761931
## 35	35	484.9241	0.09224176	382.3539	6.774730	0.008118308	5.440576
## 36	36	485.5814	0.08987890	382.9712	7.212023	0.007422832	5.995426
## 37	37	486.0300	0.08808845	383.2211	6.837801	0.010997274	5.192937
## 38	38	484.9160	0.09230702	382.4265	6.825933	0.007955897	5.377885
## 39	39	484.9383	0.09224119	382.3979	6.845493	0.007904905	5.454267
## 40	40	485.8630	0.08878243	383.5378	6.872692	0.011517395	5.523493
## 41	41	485.9552	0.08844932	383.4655	7.191351	0.009082226	6.269037
## 42	42	485.2149	0.09115320	382.6757	6.963175	0.012253751	6.061513
## 43	43	486.5938	0.08609878	383.9024	7.355962	0.011816991	5.926987
## 44	44	486.0500	0.08788994	383.2899	6.729550	0.014460579	6.331175
## 45	45	485.2134	0.09127787	382.7016	7.323259	0.007956270	6.085131
## 46	46	484.5167	0.09383109	382.1166	6.942287	0.008546818	5.593084
## 47	47	484.5038	0.09388230	382.1055	6.990469	0.008560861	5.558944
## 48	48	485.0651	0.09173109	382.5320	7.219896	0.012584797	6.109789
## 49	49	485.0128	0.09189682	382.6478	7.080085	0.011020556	5.889433

```

## 50    50 484.3602 0.09439160 382.0224 6.910623 0.008034671 5.447121
## 51    51 486.1394 0.08778257 383.4617 7.319588 0.010242351 5.798907
## 52    52 484.9639 0.09212601 382.3885 7.245413 0.012563254 6.156724
## 53    53 485.4492 0.09022959 382.8500 7.734203 0.013992719 6.413011
## 54    54 484.3753 0.09436407 381.9395 6.939935 0.008327146 5.507337
## 55    55 485.6243 0.08969806 383.4108 8.302065 0.012885651 6.213404
## 56    56 485.7829 0.08907215 383.0307 7.249422 0.013704622 5.698266
## 57    57 485.0275 0.09184288 382.3846 6.224439 0.010343737 5.270547
## 58    58 485.1327 0.09150509 382.6997 7.200257 0.011377723 6.001603
## 59    59 484.5094 0.09388720 382.0328 6.983424 0.008495922 5.488997
## 60    60 484.4635 0.09404178 382.0172 6.989106 0.008418094 5.543709
## 61    61 485.8284 0.08901398 383.1295 7.245686 0.009261405 5.346138
## 62    62 484.4185 0.09419559 381.9301 6.843253 0.007934679 5.453687
## 63    63 485.7051 0.08952791 383.1450 8.262594 0.013823307 6.065548
## 64    64 484.8459 0.09266517 382.6583 6.912692 0.007861401 5.114431
## 65    65 484.4504 0.09410302 381.9950 6.989763 0.008148227 5.528943
## 66    66 484.8090 0.09282277 382.2367 7.142013 0.006959975 5.733611
## 67    67 484.7525 0.09291981 382.3435 7.055800 0.009557842 5.813275
## 68    68 484.8394 0.09268183 382.6223 6.865309 0.007648348 5.070152
## 69    69 484.4164 0.09424034 381.9232 6.929491 0.008309412 5.464251
## 70    70 484.4425 0.09414726 381.9606 6.955284 0.008267171 5.503048
## 71    71 484.4412 0.09415282 381.9429 6.925623 0.008284369 5.484394
## 72    72 485.0356 0.09199839 382.4423 6.949698 0.010073786 5.003972
## 73    73 484.6429 0.09343302 382.1847 6.978077 0.007138629 5.633889
## 74    74 484.8997 0.09238260 382.3506 7.244394 0.010455659 5.947993
## 75    75 484.8476 0.09261320 382.3473 6.529400 0.007125989 5.394460
## 76    76 484.6673 0.09330889 382.0695 6.914579 0.009546669 5.697471
## 77    77 485.6000 0.08975523 383.0005 6.823683 0.008540653 4.974890
## 78    78 484.6858 0.09326734 382.1116 6.964984 0.009516074 5.746210
## 79    79 484.4007 0.09429229 381.8762 6.907100 0.008032206 5.516110
## 80    80 484.6085 0.09352508 381.9307 6.821025 0.008089574 5.459625
## 81    81 484.2568 0.09478787 381.8475 6.902679 0.007690307 5.499354
## 82    82 484.4098 0.09425764 381.8861 6.903003 0.007888842 5.546070
## 83    83 484.2700 0.09473658 381.8822 6.883396 0.007640213 5.491572
## 84    84 484.4261 0.09420021 381.9084 6.864125 0.007925460 5.520485
## 85    85 484.3632 0.09441468 381.8959 6.912037 0.007895170 5.541044
## 86    86 484.4332 0.09418195 381.9478 6.908887 0.008101133 5.564170
## 87    87 484.4371 0.09415600 381.9377 6.884567 0.007823598 5.501084
## 88    88 484.4495 0.09412242 381.9593 6.911859 0.007963677 5.559307
## 89    89 484.4277 0.09419129 381.9343 6.853479 0.007932078 5.490798
## 90    90 484.4206 0.09421170 381.9414 6.891539 0.007845893 5.487461
## 91    91 484.3937 0.09431767 381.9294 6.905049 0.007951948 5.492286
## 92    92 484.4090 0.09426673 381.9354 6.908262 0.008042619 5.573192
## 93    93 484.4101 0.09426115 381.9070 6.946429 0.008019806 5.544403
## 94    94 484.3897 0.09433461 381.9036 6.906962 0.007942480 5.519409
## 95    95 484.3897 0.09433461 381.9036 6.906962 0.007942480 5.519409

```

```
step_wise_model$bestTune
```

```

##      nvmax
## 81      81

```

```
# look at the coefficients of the best tuned model with 81 variables
coef(step_wise_model$finalModel, 81)
```

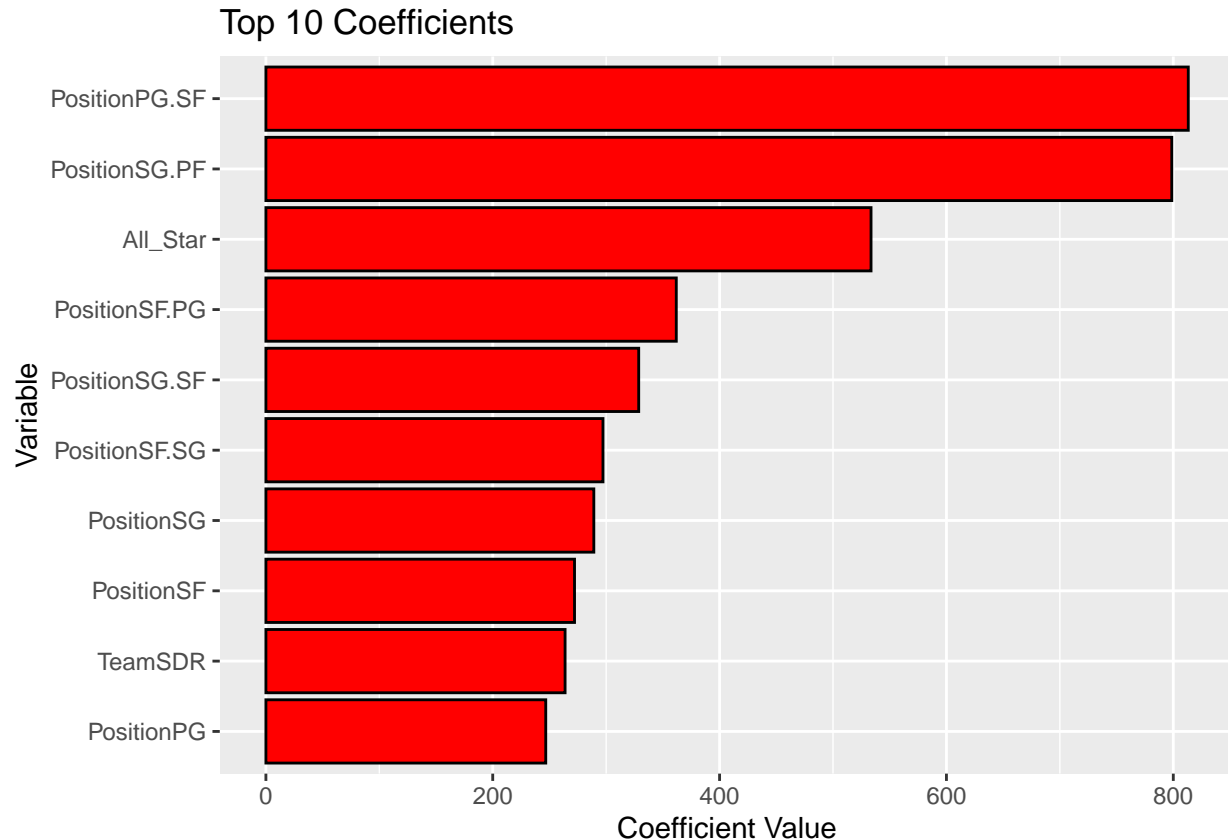
```
##      (Intercept)      height      weight      Age      All_Star
##      -501.387499      4.697110     -1.915173      1.621616     533.514839
##      TeamBAL      TeamBLB      TeamBOS      TeamBRK      TeamBUF
##      174.625306     -51.523309      44.017985     -185.554636      78.940618
##      TeamCHA      TeamCHH      TeamCHI      TeamCHO      TeamCHP
##      -122.983238     -47.391332     -17.167235     -87.535478      226.214180
##      TeamCHZ      TeamCIN      TeamCLE      TeamDAL      TeamDEN
##      96.712173      246.265185     -43.996695     -80.903992      21.118654
##      TeamDET      TeamDNN      TeamGSW      TeamKCK      TeamKCO
##      28.426896      109.293039     -11.883710      105.811260      130.470640
##      TeamLAC      TeamLAL      TeamMEM      TeamMIA      TeamMIN
##      -95.928525      100.069874     -147.950763     -91.276915     -72.826848
##      TeamMLH      TeamMNL      TeamNJN      TeamNOH      TeamNOJ
##      -157.021109      95.741644     -77.200656     -142.930054      37.375230
##      TeamNOK      TeamNOP      TeamNYK      TeamNYN      TeamOKC
##      -148.648454     -162.411909      18.667647     -140.493239     -60.527435
##      TeamORL      TeamPHI      TeamPHO      TeamPHW      TeamPOR
##      -50.667157      12.880604      21.792300      129.835745      14.777697
##      TeamROC      TeamSAC      TeamSAS      TeamSDC      TeamSDR
##      50.348134      -49.998015     -45.710062      64.413153      263.759778
##      TeamSEA      TeamSFW      TeamSHE      TeamSTB      TeamSTL
##      67.201101      237.287996      66.570249      170.419331      210.479324
##      TeamSYR      TeamTOR      TeamTOT      TeamTRI      TeamVAN
##      150.016038     -142.253415     -111.669845     -83.949338     -125.248867
##      TeamWAS      TeamWAT      TeamWSB      PositionC      PositionC.PF
##      -140.386283     -134.946396      28.967624      138.960940      121.550697
##      PositionF      PositionF.G      PositionG      PositionPF      PositionPF.C
##      -201.786045     -123.718153     -109.744151      199.717298      68.096217
##      PositionPF.SF      PositionPG      PositionPG.SF      PositionPG.SG      PositionSF
##      233.600467      246.719202      813.215296      236.273385      272.090001
##      PositionSF.PF      PositionSF.PG      PositionSF.SG      PositionSG      PositionSG.PF
##      240.055875      361.796939      297.171836      289.108783      798.662326
##      PositionSG.PG      PositionSG.SF
##      228.439915      328.690076
```

```
# predict on the test set
y_pred_MLR3 = predict.train(step_wise_model, test_set[-1], type="raw")
# compute RMSE
rmse_MLR3 = rmse(test_set$Points, y_pred_MLR3)

# plot the 10 highest coefficient magnitude variables
index_of_top_10 = c(which(abs(coef(step_wise_model$finalModel, 81))>246.3))[-1]
var_names = c('PositionPG.SF', 'PositionSG.PF', 'All_Star', 'PositionSF.PG', 'PositionSG.SF',
              'PositionSF.SG', 'PositionSG', 'PositionSF', 'TeamSDR', 'PositionPG')
coef_data = data.frame(coef(step_wise_model$finalModel, 81)[index_of_top_10])
colnames(coef_data) = c("Coef")
coef_data = coef_data %>% arrange(desc(Coef)) %>% filter(Coef>246.3)
coef_data = data.frame(cbind(var_names, coef_data))
colnames(coef_data)= c('Variable', 'Coefficient')
```



```
ggplot(coef_data, aes(x=Coefficient, y = reorder(Variable, Coefficient))) +
  geom_bar(stat = 'identity', color = "black", fill = "red") +
  labs(title = "Top 10 Coefficients", y = "Variable",
        x = "Coefficient Value")
```



This plot shows the top 10 variables of the multiple linear regression model that have the highest magnitude coefficient. The Point Guard - Shooting Forward and Shooting Guard - Power Forward positions as well as whether or not the player is an all-star contribute the highest increases in points per season, holding all other variables constant. This is fairly intuitive because those 2 positions are the ones that are taking the most shots at the net in game and have the most scoring opportunities relative to other positions. Being an all-star means that the player is more talented than the average NBA player, so these all-star players by default would be expected to be scoring more points per season than a normal player. It could also be the other way around, where all-star players are picked as all-stars because of their high scoring nature. Regardless, an all-star that is either a PG-SF or SG-PF position would have the highest season scoring potential based on this regression model.


KNN Regression

The next model I am going to test out for this dataset in predicting the number of points scored in a season by a player is KNN regression. To do this I fitted the KNN model using the scaled training and test sets, since this algorithm utilizes euclidean distance metrics in its analysis. The one hyperparameter that you must tune for KNN models is the k value. To figure out the best k-value to use in the model to get the lowest RMSE, I loop through the k-values 1-20 and build a KNN model using each k-value. I then predict the test set points using this newly created model and plot the RMSE for each k-value. You then want to pick the k-value that minimizes the RMSE, which appears to be at $k = 19$ in this case.

```

set.seed(123)
a <- NULL
for (i in 1:20) {
  model <- kknk(Points ~ .,scaled_training_set,scaled_test_set,k=i,kernel="rectangular")
  a[i] <- rmse(scaled_test_set$Points,predict(model))
}
plot(1:20,a,xlab="k",ylab="RMSE", ylim =c(400,550))

```



```

modelFinal <- kknk(Points ~ .,scaled_training_set,scaled_test_set,k=19,kernel="rectangular")
rmse_KNN <- rmse(scaled_test_set$Points,predict(modelFinal))

```

Random Forest

The next model I tested out was a random forest ensemble regression model. A random forest is basically a “forest” of decision trees where it takes the average of the result returned from all of the decision trees. The advantage of the Random Tree Model is that it is more robust than a single decision tree since it takes the average of many trees that are all built separately, allowing it to predict new values more accurately. The disadvantage of the random forest is that you must choose the number of trees in the forest as a hyperparameter. You want to use enough trees to allow the model to make accurate predictions, but you also don’t want the model to overfit to the data. Like a decision tree, random forests do not need normalized values, so I used the original training and test sets to build the model.

Once the model is built, I found the optimal hyperparameter for the number of trees in the forest by testing out different benchmark values as shown below.

- 25 trees: RMSE = 392.8573
- 50 trees: RMSE = 390.6675
- 75 trees: RMSE = 389.6562
- 100 trees: RMSE = 388.7601
- 300 trees: RMSE = 387.6793

From these findings, I chose to use 50 trees in the final model. After 50 trees, there is diminishing improvements in the RMSE, but the number of trees is still low enough to avoid overfitting to the data. 50 is the right balance between accurate predictions (minimizing RMSE) and avoiding overfitting in my opinion.

```
set.seed(123)
RFModel <- randomForest(x = training_set[-1], y = training_set$Points, ntree = 50)
y_pred_RF <- predict(RFModel, test_set)
rmse_RF <- rmse(test_set$Points, y_pred_RF)
```

SVR (Support Vector Regression)

Here we fit the SVR (Support vector regression) model to the dataset using the radial (non-linear) kernel and the eps-regression type as hyperparameters. The SVR algorithm is very effective with outliers and on non-linear problems, so I wanted to see how it compares to the other models built earlier on in the analysis. I also used the scaled training and test sets as the datasets because the SVR algorithm uses distance metrics in its computations, so the variables must be on the same scale. From the report of the RMSE, this model doesn't seem to be as effective as the previous models, so I am going to try a different kernel.

```
regressor_SVR = svm(formula = Points ~ ., data = scaled_training_set, type = 'eps-regression',
                    kernel = 'radial')
y_pred_SVR = predict(regressor_SVR, scaled_test_set)
rmse_NLSVR <- rmse(scaled_test_set$Points, y_pred_SVR)
```

Here I fit the same SVR model using a linear kernel to see if that improves the RMSE. Unfortunately the linear kernel did worse than the radial (non-linear) kernel. We can also try using the polynomial, gaussian, and sigmoid kernels to see if these can give us better results. The RMSE of these other kernel types are shown below.

- linear: 433.3105
- radial: 427
- polynomial: 433.6522
- sigmoid: 961.5413

```
regressor_SVR_Linear = svm(formula = Points ~ ., data = scaled_training_set, type = 'eps-regression',
                           kernel = 'linear')
y_pred_SVR_Linear = predict(regressor_SVR_Linear, scaled_test_set)
rmse_LSVR <- rmse(scaled_test_set$Points, y_pred_SVR_Linear)
```

Step 3: Evaluating the Performance and Describing Findings

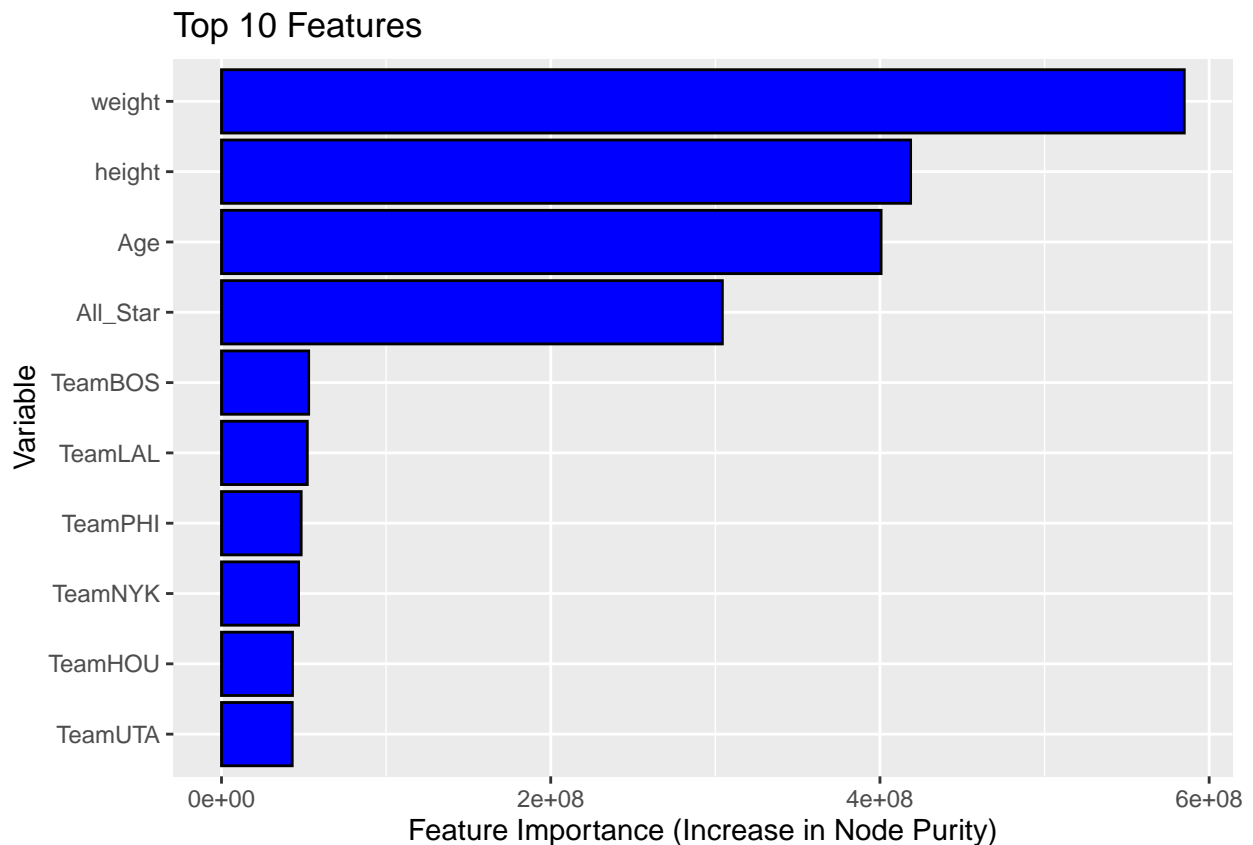
This chart shows the model name and respective RMSE value for each of the models that I built in this analysis. To choose the best model for use in predicting a player's points scored in a future season given their data, I would personally pick the Random Forest Regression model. This model got the smallest RMSE with only 390.6675 and is still not overfitted to the data because of the low number of trees used. Ensemble methods are very powerful for regression, and because of this I think the Random Forest model is the best

overall candidate for the final prediction model. I am confident in this model's ability to be able to accurately forecast how many points an NBA player is expected to score during a season, as long as it has access to all the data inputted into the model.

```
##           Model      RMSE
## 1 Multiple Linear Regression 426.4662
## 2           KNN Regression 423.1054
## 3           Random Forest 390.6675
## 4       Non-Linear SVR 427.3704
## 5           Linear SVR 433.3105
```

```
feature_importances = importance(RFModel)
var_names = colnames(training_set[-1])
feat_imp_plot_data = data.frame(var_names, feature_importances) %>% arrange(desc(IncNodePurity)) %>%
  filter(IncNodePurity>=42951190)

ggplot(feat_imp_plot_data, aes(x = IncNodePurity, y=reorder(var_names, IncNodePurity))) +
  geom_bar(stat = 'identity', color = "black", fill = "blue") +
  labs(title = "Top 10 Features", y = "Variable",
       x = "Feature Importance (Increase in Node Purity)")
```



This plot shows the feature importance values (in terms of increase in node purity) for the top 10 most important variables from the optimal model. It looks that weight, height, age, and all-star status are all very predictive of the amount of points a given player will score in a season. The teams that the players are on are still predictive of season points, but not nearly to the scale of the former variables mentioned. It seems that physical build and age of players is more important than I would have guessed and actually is more

important than all-star status as to scoring ability of the player. This somewhat surprised me, as I would think that all-star players would be top performers and that they would be able to overpower this physical traits, but it also could be a problem with imbalanced data. There are a lot less all-star players than non all-star players in the dataset, which could have caused this feature to become undervalued in the model for predicting season points. It makes sense that the team the player is on has a small role in the number of points a player scores. If a given player is on the court with extremely talented or untalented teammates, it will affect the players ability to get the ball in shooting positions and therefore affect the points scored throughout the season.